

# VIRTUALIZATION: THE CPU

Andrea Arpaci-Dusseau  
CS 537, Fall 2019

# ADMINISTRIVIA

- Project 1 is out! Due Monday, 9/16 before midnight
  - Solo, but later ones will involve project partners
  - Handin directories and test cases available
- Discussion sections on Wednesday
  - Can attend others if room (last two have most space)
- Sign up for Piazza
- Lecture recordings
- Still on waitlist? Sign attendance sheet at end of lecture
- Lecture ends at 12:15pm
- Microphone: Let me know if you can't hear

# AGENDA / OUTCOMES

## Abstraction for CPU

What is a Process? What is its lifecycle?

## Mechanism

How does a process interact with the OS?

How does the OS switch between processes?

# REVIEW

What is an Operating System?

- Software that converts hardware into a useful form for applications

What abstraction does the OS provide for the CPU?

- Process or thread

For memory?

- Virtual address space

What are some advantages of OS providing resource management?

- Protect applications from one another
- Provide fair and efficient access to resources (cost, time, energy)

# VIRTUALIZING THE CPU

High-level Goal:

Give each “process” impression it alone is actively using CPU

Resources can be shared in **time** and/or **space**

Assume single uniprocessor

- Time-sharing (multi-processors: advanced issue with space-sharing)

Memory?

- Space-sharing (later)

Disk?

- Space-sharing (later)

**ABSTRACTION: PROCESS**

# PROGRAM VS PROCESS

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"
```

```
int main(int argc, char *argv[]) {
    char *str = argv[1];
```

```
    while (1) {
        printf("%s\n", str);
        Spin(1);
```

```
    }
    return 0;
```

```
}
```

Static



Program

Running



Process

# WHAT IS A PROCESS?

Stream of executing instructions and their “context” in address space

Instruction  
Pointer



Stack pointer

```
pushq    %rbp
movq     %rsp, %rbp
subq     $32, %rsp
movl     $0, -4(%rbp)
movl     %edi, -8(%rbp)
movq     %rsi, -16(%rbp)
cmpl     $2, -8(%rbp)
je       LBB0_2
```


Registers  
Memory addrs



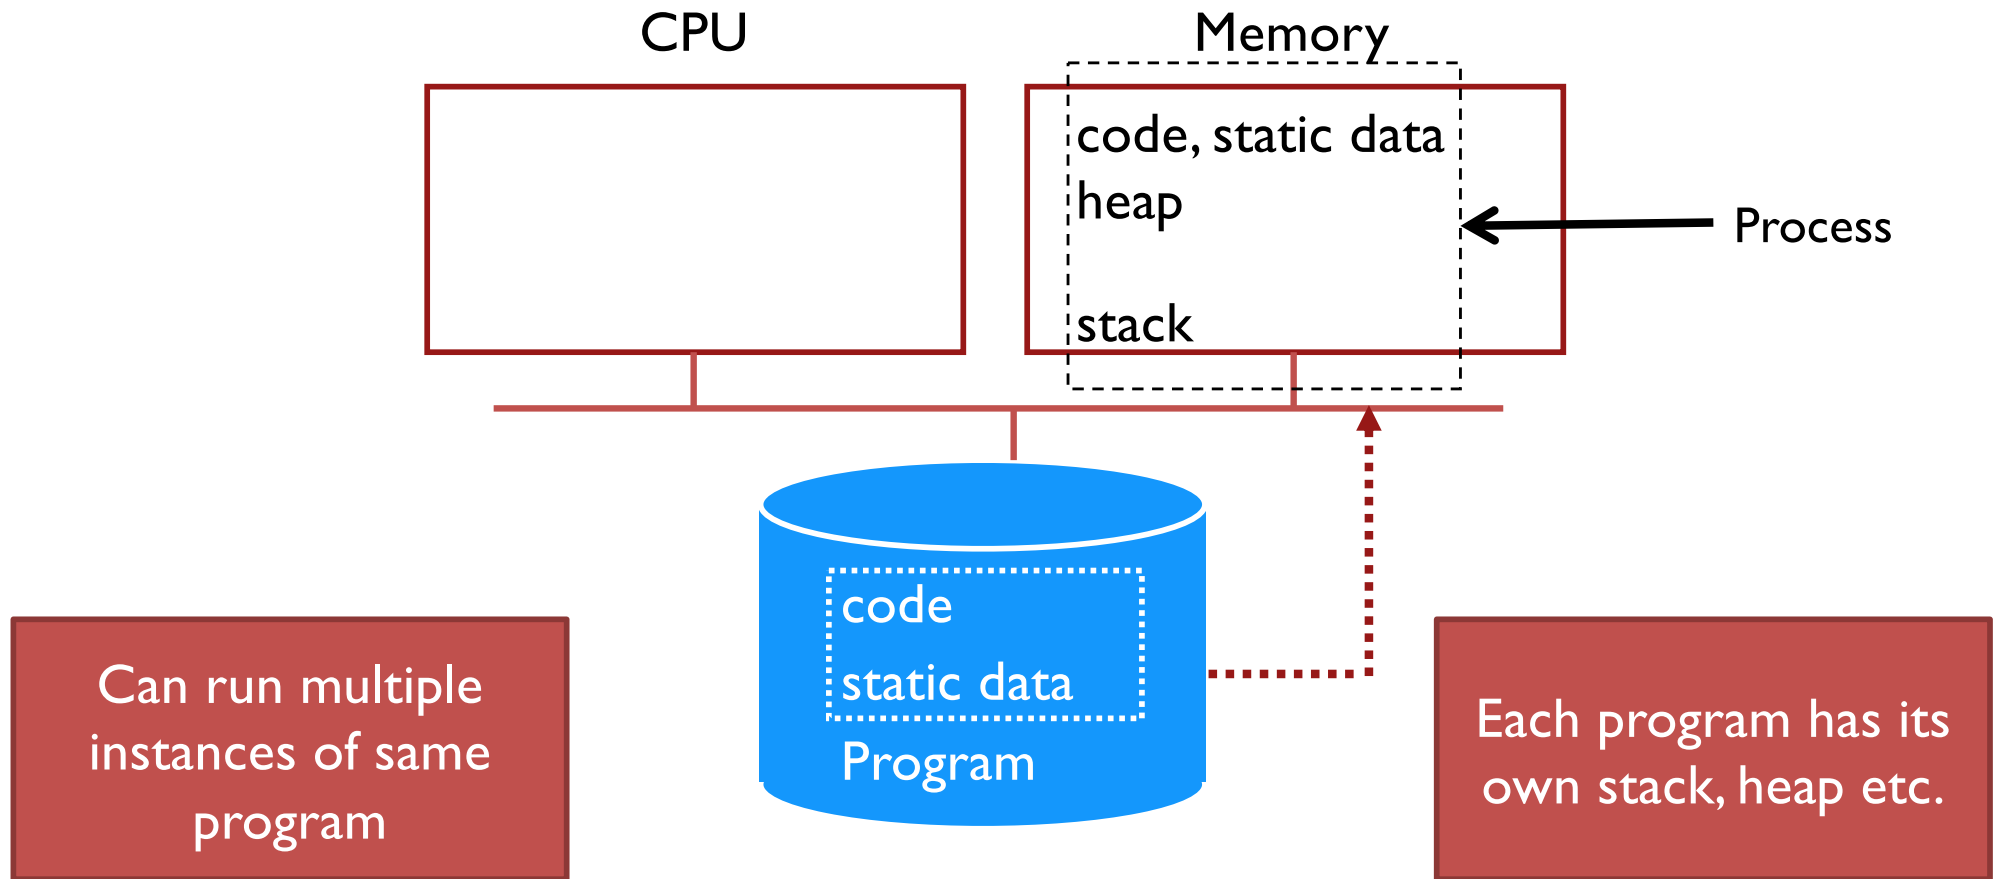
File descriptors



# PROCESS CREATION



# PROCESS CREATION



# PROCESS VS THREAD DEMO

- Two **processes** examining same memory address see **different** values (i.e., different contents)
  - Different isolated address spaces
- Two **threads** examining memory address see **same** value (i.e., same contents)
  - Share same address space

# PROCESS VS THREAD

Threads: “Lightweight process”


Execution streams that share an address space

Can directly read / write same memory

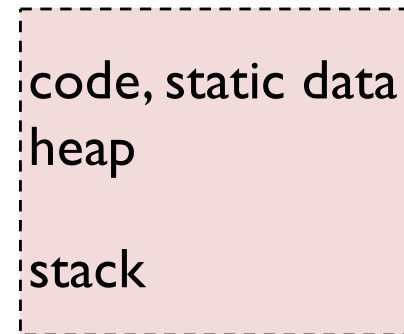
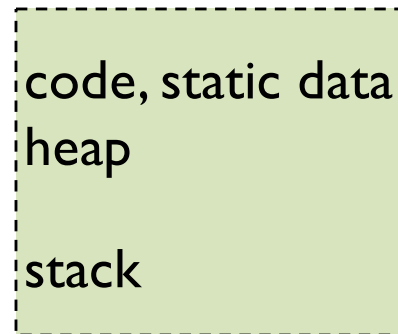
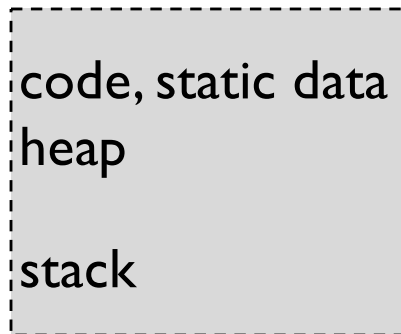
Can have multiple threads within a single process

**WHY DO WE NEED PROCESSES ?**


# SHARING CPU



# TIME SHARING




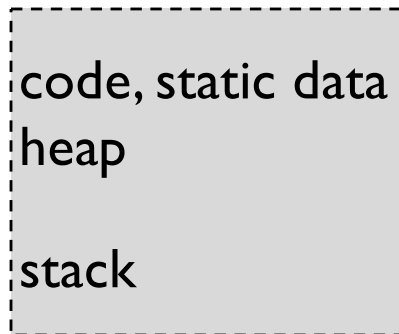


Context is loaded into CPU



CPU

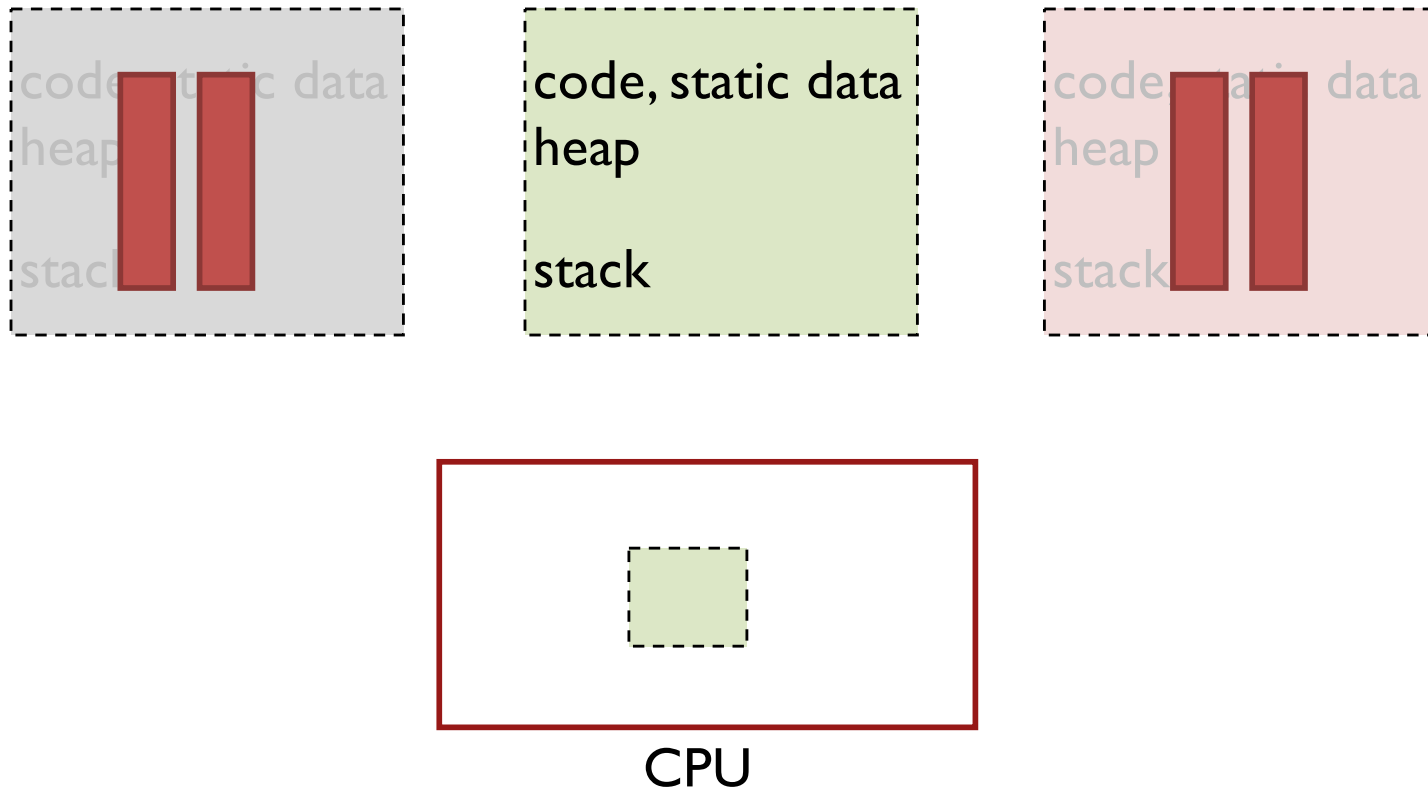
# TIME SHARING



CPU



# TIME SHARING



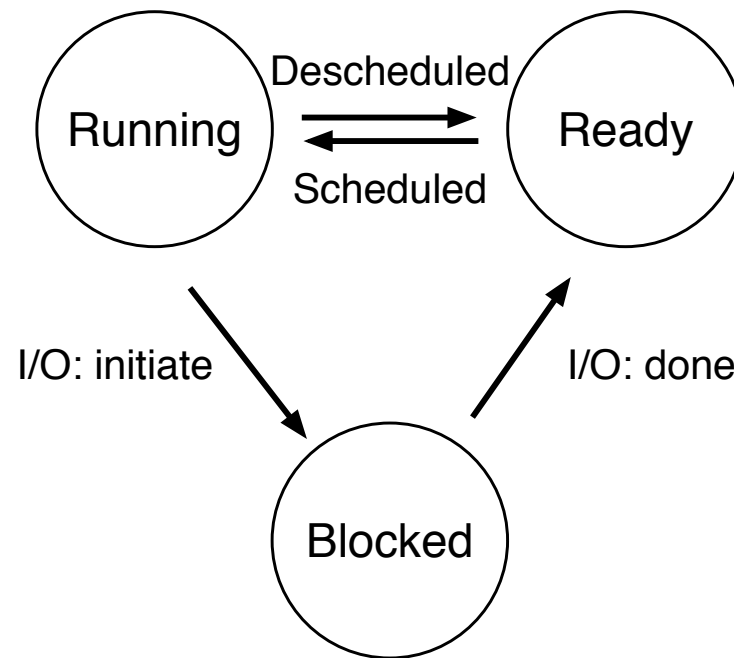
# WHAT TO DO WITH PROCESSES THAT ARE NOT RUNNING ?

OS Scheduler

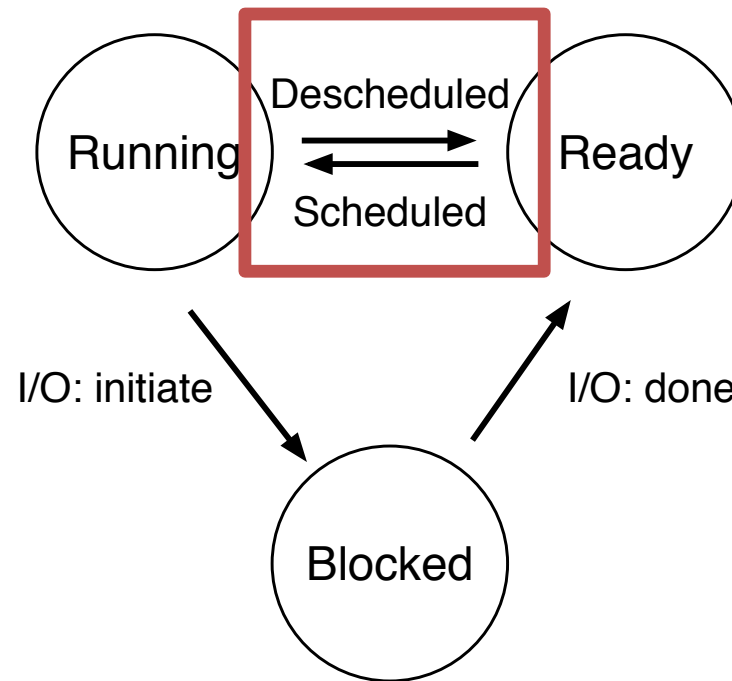
Save **context** when process is paused

Restore context on resumption

# STATE TRANSITIONS



# STATE TRANSITIONS



# ASIDE: OSTEP HOMEWORKS!

- Optional homeworks corresponding to each chapter in book
- Little simulators to help you understand
- Can generate problems and solutions!

<http://pages.cs.wisc.edu/~remzi/OSTEP/Homework/homework.html>

# PROCESS HW

Run `./process_run.py -l 2:100,2:0`

# QUIZ

$\geq$  ./process-run.py -l 3:50,3:40

Process 0

io

io

cpu

Process 1

cpu

io

io

# CPU TIME SHARING

Mechanism goals: Be able to run processes

Efficiency: Time sharing should not add overhead

Control: OS should be able to intervene when required

Policy goals: Pick the “best” process to schedule

Reschedule process for fairness? efficiency ?

Separate mechanism from policy for clean OS design

How to have efficient mechanism??



# EFFICIENT EXECUTION

Simple answer !?: **Direct Execution**

Allow user process to run directly

Create process and transfer control to main()

Challenges

- 1) What if the process wants to do something restricted ? Access disk ?
- 2) What if the process runs forever ? Buggy ? Malicious ?

Solution: **Limited Direct Execution (LDE)**

# CHALLENGE 1: RESTRICTED OPS

How can we ensure user process can't harm others?

Solution: privilege levels supported by hardware (bit of status)

User processes run in user mode (restricted mode)

OS runs in kernel mode (not restricted)

How can process access devices?

**System calls** (function call implemented by OS)

**SYSTEM CALL**

# SYSTEM CALL



P wants to call read()

# SYSTEM CALL



P can only see its own memory because of **user mode**  
(other areas, including kernel, are hidden)

# SYSTEM CALL




P wants to call `read()` but no way to call it directly

# SYSTEM CALL



```
movl $6, %eax;    int $64
```

# SYSTEM CALL




`movl $6, %eax;`

`int $64`

Trap table  
index




# SYSTEM CALL



Syscall table  
index

`movl $6, %eax;      int $64`

# SYSTEM CALL




Syscall table  
index

`movl $6, %eax;`

`int $64`

Trap table  
index


# SYSTEM CALL



```
movl $6, %eax;    int $64
```

Follow entries to correct system call code

# SYSTEM CALL



Kernel can access user memory to fill in user buffer  
return-from-trap at end to return to Process P

# SYSCALL SUMMMARY

Separate user-mode from kernel mode for security

Syscall: call kernel mode functions

- Transfer from user-mode to kernel-mode (trap)

- Return from kernel-mode to user-mode (return-from-trap)

# 5 MINUTE BREAK!

Talk with at least two neighbors

- What has been your favorite course in CS?  
What did you like most about it?
- Favorite course outside of CS? Why?

# REPEAT: EFFICIENT EXECUTION

Simple answer !?: **Direct Execution**

Allow user process to run directly

Create process and transfer control to main()

Challenges

- 1) What if the process wants to do something restricted ? Access disk ?
- 2) What if the process runs forever ? Buggy ? Malicious ?

Solution: **Limited Direct Execution (LDE)**

# CHALLENGE 2: HOW TO TAKE CPU AWAY

## Policy

To decide which process to schedule when

Decision-maker to optimize some workload performance metric

## Mechanism

To switch between processes

Low-level code that implements the decision


Separation of policy and mechanism: Recurring theme in OS



# DISPATCH MECHANISM

OS runs **dispatch loop**

```
while (1) {  
    run process A for some time-slice  
    stop process A and save its context  
    load context of another process B  
}
```

A red bracket on the right side of the code block groups the lines "stop process A and save its context" and "load context of another process B". To the right of this bracket is the text "Context-switch".

Context-switch


Question 1: How does dispatcher gain control?

Question 2: What must be saved and restored?

# HOW DOES DISPATCHER GET CONTROL?

Option 1: **Cooperative Multi-tasking**: Trust process to relinquish CPU through traps

- Examples: System call, page fault (access page not in main memory), or error (illegal instruction or divide by zero)
- Provide special `yield()` system call



# PROBLEMS WITH COOPERATIVE ?

Disadvantages: Processes can **misbehave**

By avoiding all traps and performing no I/O, can take over entire machine

Only solution: Reboot!

Not performed in modern operating systems

# TIMER-BASED INTERRUPTS

Option 2: Timer-based Multi-tasking (True multi-tasking)

Guarantee OS can obtain control periodically

Enter OS by enabling periodic alarm clock

Hardware generates timer interrupt (CPU or separate chip) Example: Every 10ms

User must not be able to mask timer interrupt

Operating System

Hardware

Program  
Process A

Operating System

Hardware

Program  
Process A

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler for timer

Operating System

Hardware

Program  
Process A

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler for timer

Handle the trap for timer  
Call switch() routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

## Operating System

Handle the trap for timer  
Call switch() routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

## Hardware

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler for timer

restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

## Program Process A



## Operating System

## Hardware

## Program Process A

Handle the trap for timer  
Call switch() routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

timer interrupt  
save regs(A) to k-stack(A)  
move to kernel mode  
jump to trap handler for timer

restore regs(B) from k-stack(B)  
move to user mode  
jump to B's IP

## Process B

# SUMMARY


Process: Abstraction to virtualize CPU

Use time-sharing in OS to switch between processes

Key aspects

- Use system calls to run access devices etc. from user mode

- Context-switch using interrupts for multi-tasking



**POLICY ?**  
**NEXT CLASS!**

# ADMINISTRIVIA

- Lecture ends at 12:15pm
- Project 1 is out! Due Monday, 9/16 before midnight
  - Handin directories and test cases available
- Discussion sections on Wednesday
  - Can attend others if room (last two have most space)
- Sign up for Piazza
- Still on waitlist? Sign attendance sheet at end of lecture