

# CONCURRENCY: QUEUE LOCKS AND CONDITION VARIABLES

Andrea Arpaci-Dusseau  
CS 537, Fall 2019

# ADMINISTRIVIA

- Project 4: xv6 Scheduler
  - Due Tuesday at 5:00pm (or midnight)
  - Test Cases available
  - Handin directories available
- Project 5 (xv6 Virtual Memory) available Tuesday
  - Partners strongly recommended

# AGENDA / LEARNING OUTCOMES

## Concurrency

- How to **block** instead of **spin-wait** while waiting for a lock?
- **When** should a waiting thread block vs. spin?
- How can threads enforce **ordering** across operations (**condition variables**)?
- How can **thread\_join()** be implemented?
- How can we support **producer/consumer** apps?

**RECAP**

# LOCK IMPLEMENTATION GOALS

## Correctness

- *Mutual exclusion*  
Only one thread in critical section at a time
- *Progress* (deadlock-free)  
If several simultaneous requests, must allow one to proceed
- *Bounded* (starvation-free)  
Must eventually allow each waiting thread to enter

Fairness: Each thread given lock in same order as requested

Performance: CPU is not used unnecessarily

# LOCK IMPLEMENTATION WITH XCHG

```
typedef struct __lock_t {
    int flag;
} lock_t;

void init(lock_t *lock) {
    lock->flag = 0;
}

void acquire(lock_t *lock) {
    while(xchg(&lock->flag, 1) == 1) ;
    // spin-wait (do nothing)
}

void release(lock_t *lock) {
    lock->flag = 0;
}
```

# FAIRNESS: TICKET LOCKS

Idea: reserve each thread's turn to use a lock.

Each thread spins until their turn.

Use new atomic primitive, fetch-and-add

```
int FetchAndAdd(int *ptr) {  
    int old = *ptr;  
    *ptr = old + 1;  
    return old;  
}
```

Acquire: Grab ticket; Spin while not thread's ticket != turn

Release: Advance to next turn

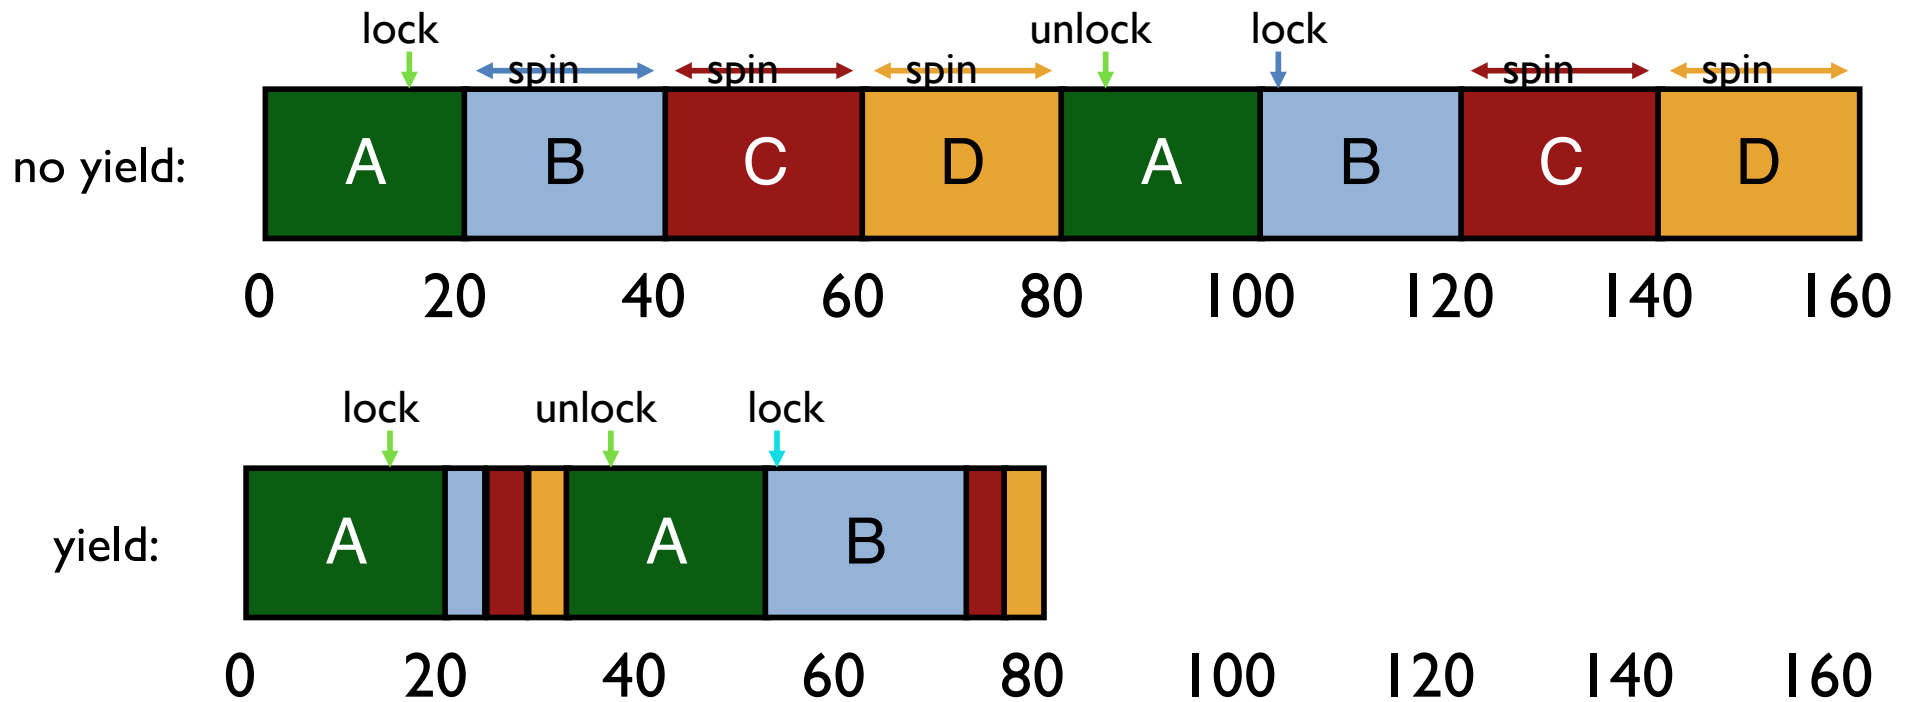
# TICKET LOCK WITH YIELD

```
typedef struct __lock_t {  
    int ticket;  
    int turn;  
}  
  
void lock_init(lock_t *lock) {  
    lock->ticket = 0;  
    lock->turn = 0;  
}  
  
void acquire(lock_t *lock) {  
    int myturn = FAA(&lock->ticket);  
    while (lock->turn != myturn)  
        yield();  
}  
  
void release(lock_t *lock) {  
    FAA(&lock->turn);  
}
```

Remember: `yield()` voluntarily relinquishes CPU for remainder of timeslice, but process remains READY



# YIELD INSTEAD OF SPIN



# SPINLOCK PERFORMANCE

Waste of CPU cycles?

Without yield:  $O(\text{threads} * \text{time\_slice})$

With yield:  $O(\text{threads} * \text{context\_switch})$

Even with yield, spinning is slow with high thread contention

Next improvement: Block and put thread on waiting queue instead of spinning

# QUEUE LOCKS

# LOCK IMPLEMENTATION: BLOCK WHEN WAITING

Remove waiting threads from scheduler ready queue

Move to BLOCKED or WAITING state

(e.g., park() and unpark(threadID))

Scheduler runs any thread that is **ready**

Good separation of concerns between lock and scheduler

RUNNABLE: A, B, C, D

RUNNING: <empty>

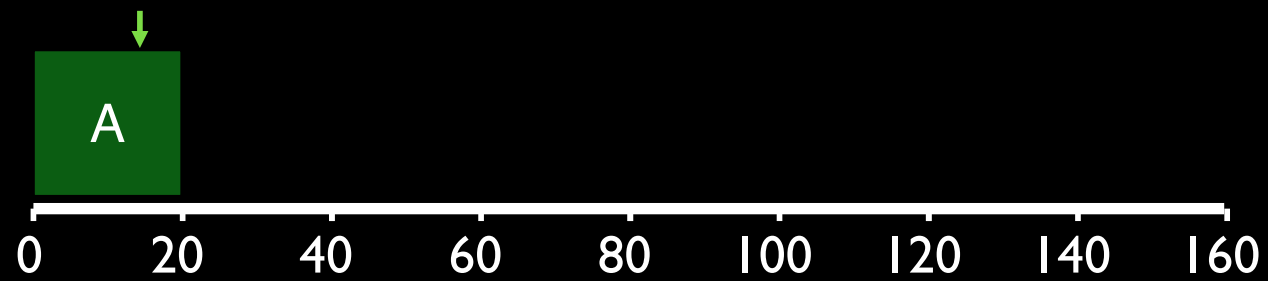
WAITING: <empty>



RUNNABLE: B, C, D

RUNNING: A

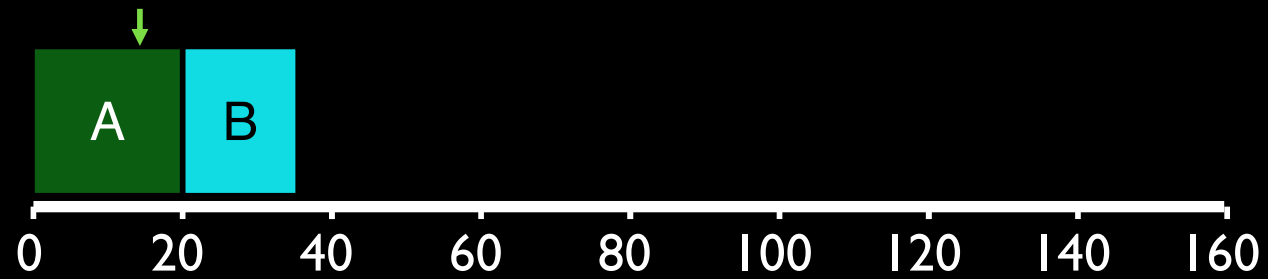
WAITING: <empty>



RUNNABLE: C, D, A

RUNNING: B

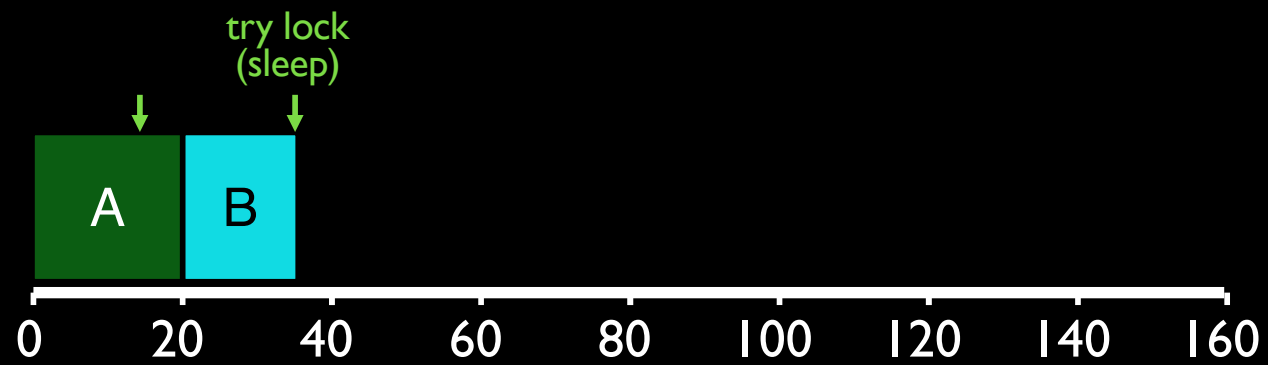
WAITING: <empty>



RUNNABLE: C, D, A

RUNNING:

WAITING: B

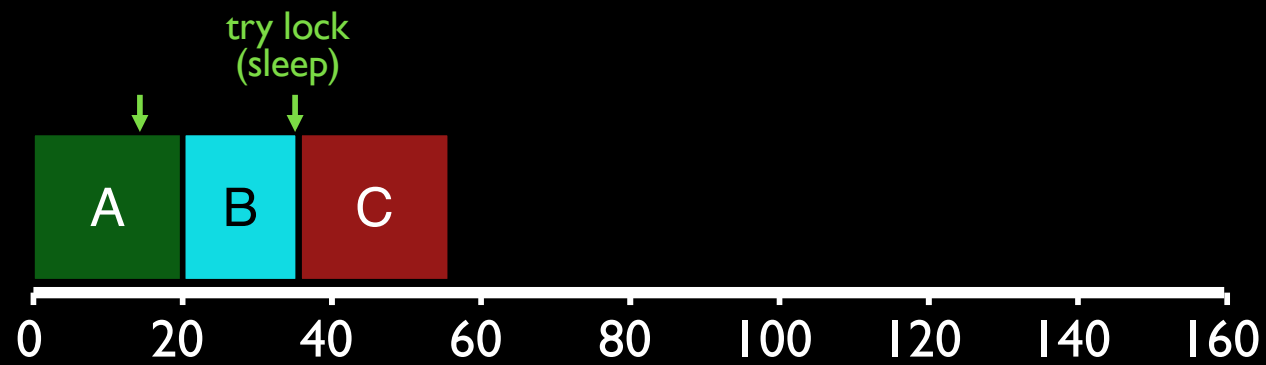




RUNNABLE: D,A

RUNNING: C

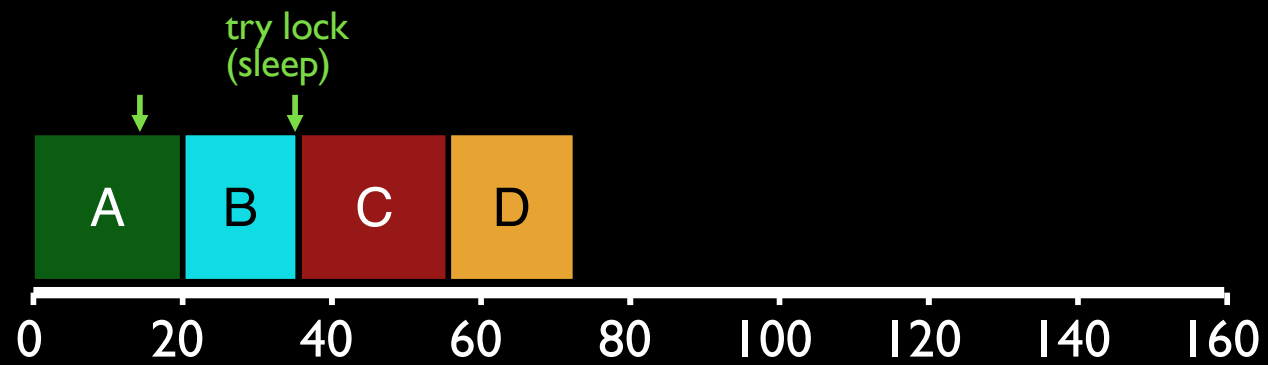
WAITING: B



RUNNABLE: A, C

RUNNING: D

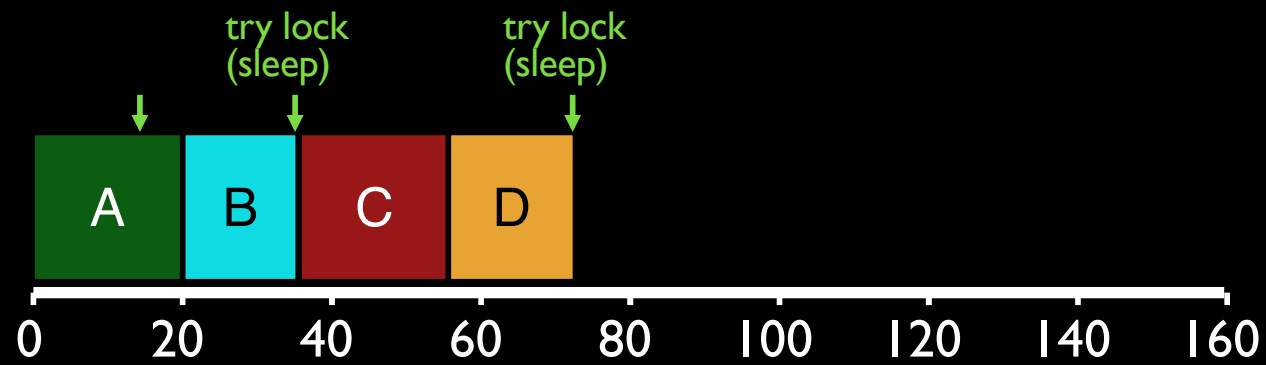
WAITING: B



RUNNABLE: A, C

RUNNING:

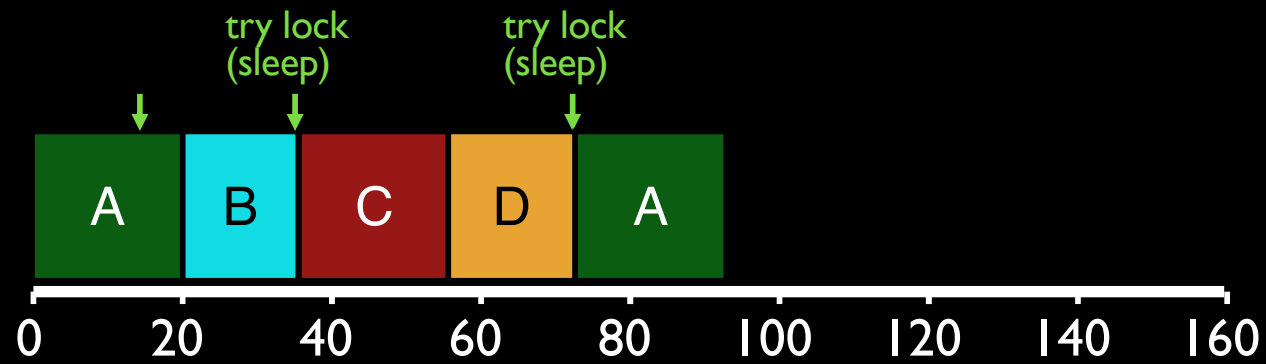
WAITING: B, D



RUNNABLE: C

RUNNING: A

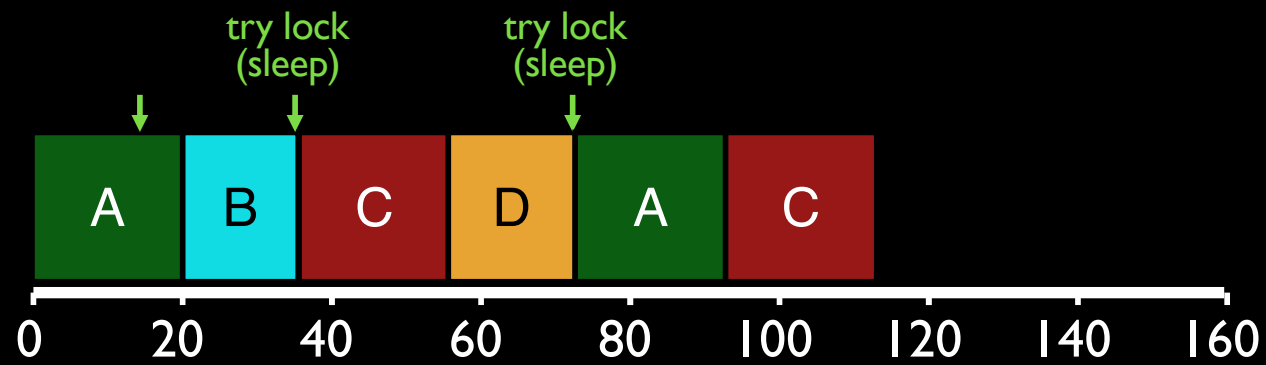
WAITING: B, D



RUNNABLE: A

RUNNING: C

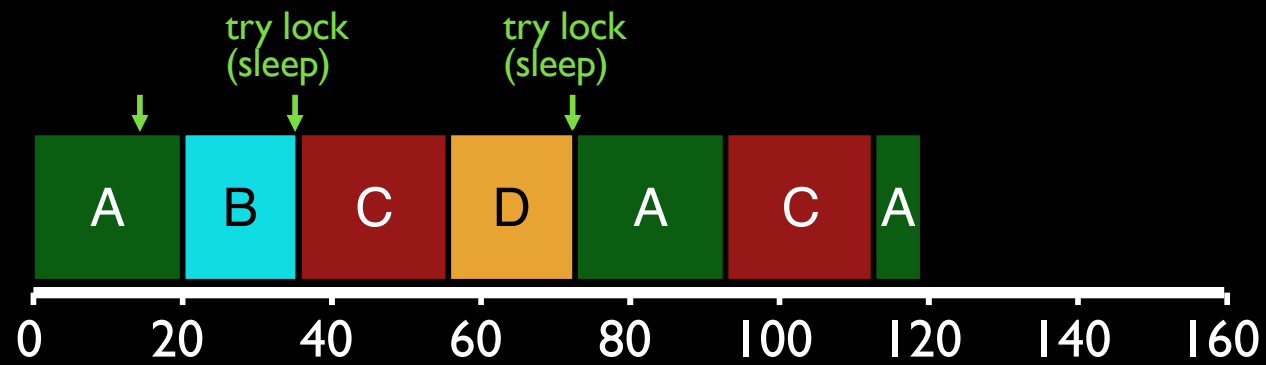
WAITING: B, D



RUNNABLE: C

RUNNING: A

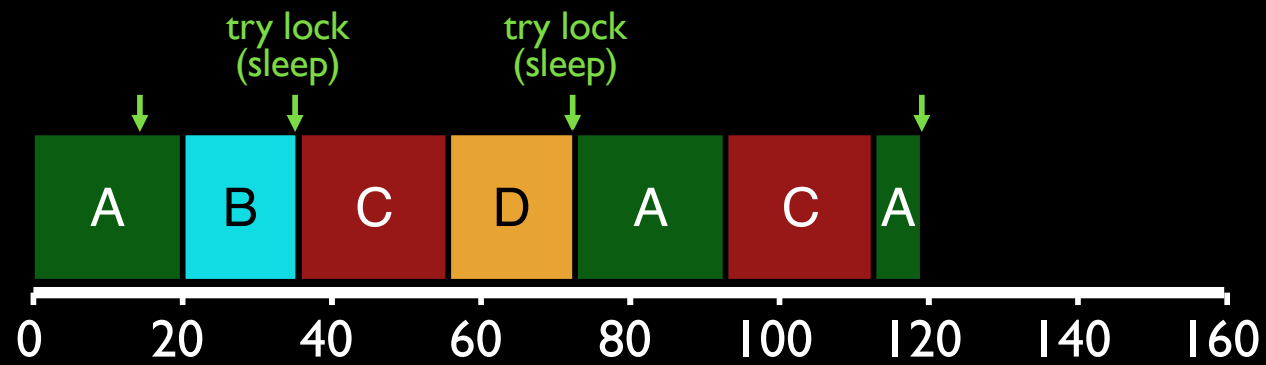
WAITING: B, D



RUNNABLE: B, C

RUNNING: A

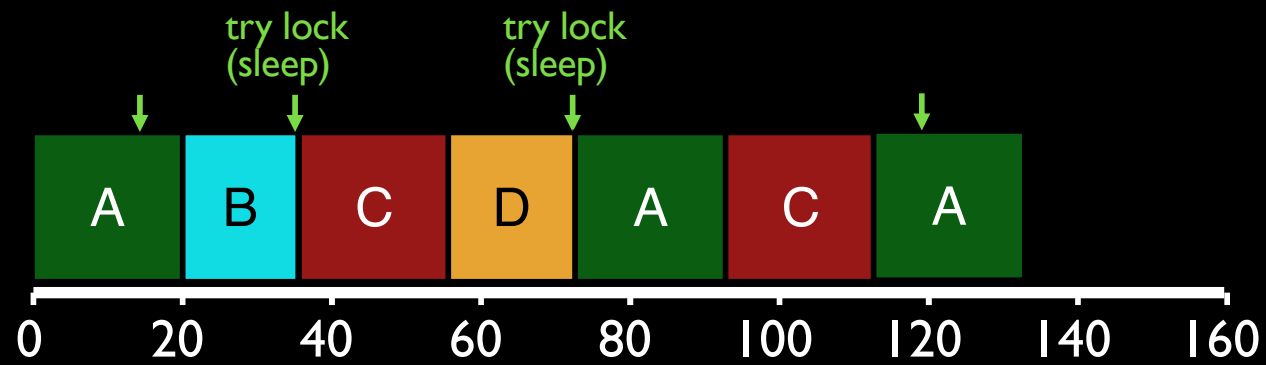
WAITING: D



RUNNABLE: B, C

RUNNING: A

WAITING: D

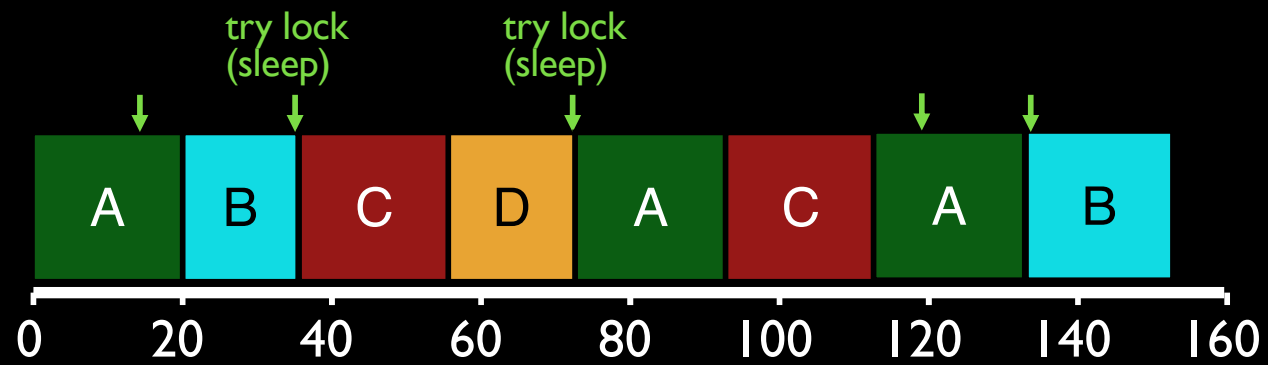




RUNNABLE: C,A

RUNNING: B

WAITING: D



# LOCK IMPLEMENTATION #1: BLOCK WHEN WAITING

```
typedef struct {  
    bool lock = false;  
    queue_t q;  
} LockT;
```

Track waiting processes on q

```
void acquire(LockT *l) {  
    if (l->lock) {  
        qadd(l->q, tid);  
        park();    // blocked  
    } else {  
        l->lock = true;  
    }  
}  
  
void release(LockT *l) {  
    if (qempty(l->q)) l->lock=false;  
    else unpark(qremove(l->q));  
}
```

# LOCK IMPLEMENTATION #2: BLOCK WHEN WAITING

```
typedef struct {  
    bool lock = false;  
    bool guard = false;  
    queue_t q;  
} LockT;
```

Add guard to lock

```
void acquire(LockT *l) {  
    while (XCHG(&l->guard, true));  
    if (l->lock) {  
        qadd(l->q, tid);  
        l->guard = false;  
        park();    // blocked  
    } else {  
        l->lock = true;  
        l->guard = false;  
    }  
}
```

```
void release(LockT *l) {  
    while (XCHG(&l->guard, true));  
    if (qempty(l->q)) l->lock=false;  
    else unpark(qremove(l->q));  
    l->guard = false;  
}
```

# LOCK IMPLEMENTATION: BLOCK WHEN WAITING

(a) Why is **guard** used?

(b) Why okay to **spin** on guard?

(c) In release(), why not set lock=false when unpark?

(d) Is there a race condition?

```
void acquire(LockT *l) {
    while (XCHG(&l->guard, true));
    if (l->lock) {
        qadd(l->q, tid);
        l->guard = false;
        park();      // blocked
    } else {
        l->lock = true;
        l->guard = false;
    }
}
```

```
void release(LockT *l) {
    while (XCHG(&l->guard, true));
    if (qempty(l->q)) l->lock=false;
    else unpark(qremove(l->q));
    l->guard = false;
}
```

# RACE CONDITION

**Thread 1** (in lock)

```
if (l->lock) {  
    qadd(l->q, tid);  
    l->guard = false;
```

```
park();    // block
```

**Thread 2** (in unlock)

```
while (TAS(&l->guard, true));  
if (qempty(l->q)) // false!!  
else unpark(qremove(l->q));  
l->guard = false;
```

Problem: Guard not held when call park()  
Unlocking thread may unpark() before other park()

# BLOCK WHEN WAITING: FINAL CORRECT LOCK

```
typedef struct {  
    bool lock = false;  
    bool guard = false;  
    queue_t q;  
} LockT;
```

setpark() fixes race condition

Park() does not block if unpark()  
occurred after setpark()

```
void acquire(LockT *l) {  
    while (TAS(&l->guard, true));  
    if (l->lock) {  
        qadd(l->q, tid);  
        setpark(); // notify of plan  
        l->guard = false;  
        park(); // unless unpark()  
    } else {  
        l->lock = true;  
        l->guard = false;  
    }  
}  
  
void release(LockT *l) {  
    while (TAS(&l->guard, true));  
    if (qempty(l->q)) l->lock=false;  
    else unpark(qremove(l->q));  
    l->guard = false;  
}
```

# REASONING ABOUT LOCKS

- When using locks, don't assume any implementation details are necessary for correctness of calling process
- Don't assume any particular ordering for which process acquires lock next
- Your application code must work correctly if any process acquires lock next

# PERFORMANCE: SPIN-WAITING VS BLOCKING

Each approach is better under different circumstances

## Uniprocessor

Waiting process is scheduled → Process holding lock isn't

Waiting process should always relinquish processor

Associate queue of waiters with each lock (as in previous implementation)

## Multiprocessor

Waiting process is scheduled → Process holding lock might be

Spin or block depends on how long,  $t$ , before lock is released

Lock released quickly → Spin-wait

Lock released slowly → Block

Quick and slow are relative to context-switch cost,  $C$



# WHEN TO SPIN-WAIT? WHEN TO BLOCK?

If know how long,  $t$ , before lock released, can determine **optimal** behavior

How much CPU time is wasted when spin-waiting?

How much wasted when block? <sup>$t$</sup>

What is the best action when  $t < C$ ? <sup>$C$</sup>

When  $t > C$ ? spin-wait

block

Problem:

Requires knowledge of future;  
too much overhead to do any special prediction

# TWO-PHASE WAITING

Theory: Bound worst-case performance;  
ratio of actual / optimal

When would worst-possible performance occur?

Spin for very long time  $t \gg C$   
Ratio:  $t/C$  (unbounded)

Algorithm: Spin-wait for time  $C$  then block  $\rightarrow$  Factor of 2 of optimal

Two cases:

$t < C$ : optimal approach spin-waits for  $t$ ; we also spin-wait  $t$

$t > C$ : optimal blocks immediately (cost of  $C$ ); we pay spin  $C$  then block (cost of  $2C$ );  
 $2C / C \rightarrow 2$ -competitive algorithm

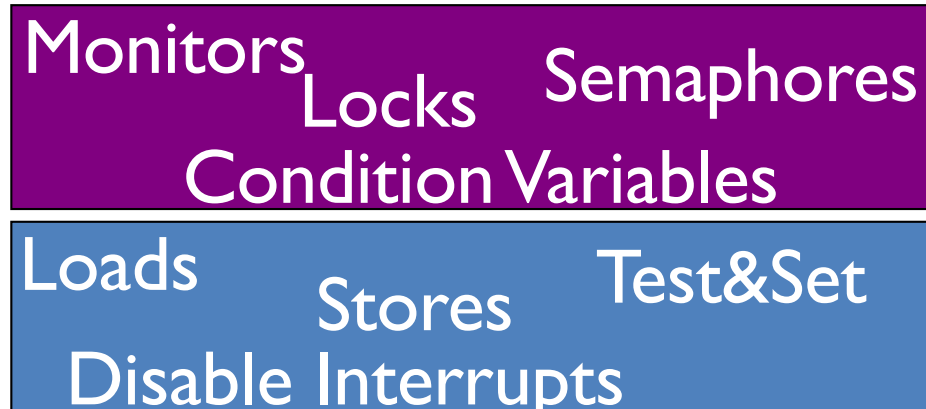
Example of **competitive analysis**

# IMPLEMENTING SYNCHRONIZATION

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



# CONDITION VARIABLES

# CONCURRENCY OBJECTIVES

**Mutual exclusion** (e.g., A and B don't run at same time)

- solved with *locks*

**Ordering** (e.g., B runs after A does something)

- solved with *condition variables* and *semaphores*

# ORDERING EXAMPLE: JOIN

```
pthread_t p1, p2;  
Pthread_create(&p1, NULL, mythread, "A");  
Pthread_create(&p2, NULL, mythread, "B");  
  
// join waits for specified thread to finish  
Pthread_join(p1, NULL);  
Pthread_join(p2, NULL);  
  
printf("main: done\n [balance: %d]\n [should: %d]\n",  
       balance, max*2);  
return 0;
```

how to implement join()?

# CONDITION VARIABLES

Condition Variable: queue of waiting threads

**B** waits for a signal on CV before running

- wait(CV, ...)

**A** sends signal to CV when time for **B** to run

- signal(CV, ...)

# CONDITION VARIABLES

**wait**(cond\_t \*cv, mutex\_t \*lock)

- assumes the specified lock is held when wait() is called
- puts caller to sleep + releases the lock (atomically)
- when awoken, reacquires lock before returning

**signal**(cond\_t \*cv)

- wake a single waiting thread (if  $\geq 1$  thread is waiting)
- if there is no waiting thread, just return, doing nothing



# JOIN IMPLEMENTATION: ATTEMPT 1

Parent:

```
void thread_join() {  
    Mutex_lock(&m);    // x  
    Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);  // z  
}
```

Child:

```
void thread_exit() {  
    Cond_signal(&c);    // a  
}
```

Example schedule:

Works!?

Parent:	x	y	z
Child:		a	

# JOIN IMPLEMENTATION: ATTEMPT 1

Parent:

```
void thread_join() {  
    Mutex_lock(&m);    // x  
    Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);  // z  
}
```

Child:

```
void thread_exit() {  
    Cond_signal(&c);    // a  
}
```

Parent waits forever!

Example broken schedule:

Parent:		x	y
Child:	a		

# CV RULE OF THUMB 1

Keep state in addition to CV's!

CV's are used to signal threads when state changes

If state is already as needed, thread doesn't wait for a signal!

# JOIN IMPLEMENTATION: ATTEMPT 2

Int done = 0; // shared between parent and child

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;                 // a  
    Cond_signal(&c);          // b  
}
```

Fixes previous broken ordering:

Parent:

w

x

y

z

Child:

a

b

# JOIN IMPLEMENTATION: ATTEMPT 2

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)            // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    done = 1;                 // a  
    Cond_signal(&c);          // b  
}
```

Parent waits forever!

Can you construct ordering that does not work?

Parent: w      x                      y

Child:                      a      b

# JOIN IMPLEMENTATION: CORRECT

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1;                // b  
    Cond_signal(&c);          // c  
    Mutex_unlock(&m);         // d  
}
```

Parent: w	x	y	z	
Child:	?	a	b	c

Use mutex to ensure no race between interacting with state and wait/signal  
Essential that mutex is released within cond wait()

# CV RULE OF THUMB 2

Modify state with mutex held (in threads calling wait and signal)

Mutex is required to ensure state does not change between testing of state and waiting on CV

# PRODUCER/CONSUMER PROBLEM



# EXAMPLE: UNIX PIPES

A pipe may have many writers and readers

Internally, there is a finite-sized, circular buffer

Writers add data to the buffer

- Writers have to wait if buffer is full

Readers remove data from the buffer

- Readers have to wait if buffer is empty

Buf:



# PRODUCER/CONSUMER PROBLEM

**Producers** generate data (like pipe writers)

**Consumers** grab data and process it (like pipe readers)

Producer/consumer problems are frequent in systems (e.g. web servers)

General strategy use condition variables to:

- make producers wait when buffers are full

- make consumers wait when buffers are empty

# PRODUCE/CONSUMER EXAMPLE

Start with easy case:

- 1 producer thread
- 1 consumer thread
- 1 shared buffer to fill/consume (max = 1)

Numfull = number of buffers currently filled

max = 1

numfull = 0

Thread 1 state:

```
void *producer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        if(numfull == max)
            Cond_wait(&cond, &m);
        do_fill();
        Cond_signal(&cond);
        Mutex_unlock(&m);
    }
}
```

Assume do\_fill() increments numfull

Thread 2 state:

```
void *consumer(void *arg) {
    while(1) {
        Mutex_lock(&m);
        if(numfull == 0)
            Cond_wait(&cond, &m);
        tmp = do_get();
        Cond_signal(&cond);
        Mutex_unlock(&m);
        do_work(tmp);
    }
}
```

Assume do\_get() decrements numfull

numfull = 0

Thread 1 state: RUNNABLE

```
void *producer(void *arg) {  
→ while (1) {  
    Mutex_lock(&m);  
    if(numfull == max)  
        Cond_wait(&cond, &m);  
    do_fill();  
    Cond_signal(&cond);  
    Mutex_unlock(&m);  
}  
}
```

Thread 2 state: RUNNING

```
void *consumer(void *arg) {  
→ while(1) {  
    Mutex_lock(&m);  
    if(numfull == 0)  
        Cond_wait(&cond, &m);  
    tmp = do_get();  
    Cond_signal(&cond);  
    Mutex_unlock(&m);  
    do_work(tmp);  
}  
}
```

numfull = 0

Thread 1 state: RUNNABLE

```
void *producer(void *arg) {  
→ while (1) {  
    Mutex_lock(&m);  
    if(numfull == max)  
        Cond_wait(&cond, &m);  
    do_fill();  
    Cond_signal(&cond);  
    Mutex_unlock(&m);  
}  
}
```

Thread 2 state: RUNNING

```
void *consumer(void *arg) {  
→ while(1) {  
    Mutex_lock(&m);  
    if(numfull == 0)  
        Cond_wait(&cond, &m);  
    tmp = do_get();  
    Cond_signal(&cond);  
    Mutex_unlock(&m);  
    do_work(tmp);  
}  
}
```

numfull = 0

Thread 1 state: RUNNABLE

```
void *producer(void *arg) {  
→ while (1) {  
    Mutex_lock(&m);  
    if(numfull == max)  
        Cond_wait(&cond, &m);  
    do_fill();  
    Cond_signal(&cond);  
    Mutex_unlock(&m);  
}  
}
```

Thread 2 state: RUNNING

```
void *consumer(void *arg) {  
→ while(1) {  
    Mutex_lock(&m);  
    if(numfull == 0)  
        Cond_wait(&cond, &m);  
    tmp = do_get();  
    Cond_signal(&cond);  
    Mutex_unlock(&m);  
    do_work(tmp);  
}  
}
```

numfull = 0

Thread 1 state: RUNNABLE

```
void *producer(void *arg) {  
→ while (1) {  
    Mutex_lock(&m);  
    if(numfull == max)  
        Cond_wait(&cond, &m);  
    do_fill();  
    Cond_signal(&cond);  
    Mutex_unlock(&m);  
}  
}
```

Thread 2 state: RUNNING

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```



numfull = 0

Thread 1 state: RUNNABLE

```
void *producer(void *arg) {  
    → while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


Consumer releases mutex in cond\_wait

Thread 2 state: BLOCKED on CV

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            → Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```


numfull = 0

Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

Thread 2 state: BLOCKED on CV




```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

Will producer be stuck waiting for mutex\_lock()?

numfull = 0


Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

No, because cond\_wait released lock


Thread 2 state: BLOCKED on CV



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

numfull = 0


Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

Numful != max


Thread 2 state: BLOCKED on CV



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

numfull = 1


Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

What happens to consumer?


Thread 2 state: BLOCKED on CV



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```


numfull = 1

Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

Thread 2 state: BLOCKED on MUTEX

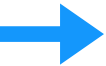


```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

Consumer must reacquire mutex to return from cond\_wait


numfull = 1

Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


Thread 2 state: RUNNABLE



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

numfull = 1

Thread 1 state: RUNNING

```
void *producer(void *arg) {  
    while (1) {  
         Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

Thread 2 state: RUNNABLE

```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

No assumptions for correctness! Could be either thread!


Who acquires lock next?

Example gives lock to producer




numfull = 1

Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```


Thread 2 state: RUNNABLE



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```


numfull = 1

Thread 1 state: RUNNING



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

Thread 2 state: RUNNABLE




```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

What important thing happens during cond\_wait()?

numfull = 1


Thread 1 state: BLOCKED on CV



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

Producer releases mutex


Thread 2 state: RUNNING, Acquires lock



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

numfull = 1


Thread 1 state: BLOCKED on CV



```
void *producer(void *arg) {  
    while (1) {  
        Mutex_lock(&m);  
        if(numfull == max)  
            Cond_wait(&cond, &m);  
        do_fill();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
    }  
}
```

And so on...

Thread 2 state: RUNNING, Acquires lock



```
void *consumer(void *arg) {  
    while(1) {  
        Mutex_lock(&m);  
        if(numfull == 0)  
            Cond_wait(&cond, &m);  
        tmp = do_get();  
        Cond_signal(&cond);  
        Mutex_unlock(&m);  
        do_work(tmp);  
    }  
}
```

# WHAT ABOUT 2 CONSUMERS?

Can you find a problematic timeline with 2 consumers (still 1 producer)?

```

void *producer(void *arg) {
    while (1) {
        Mutex_lock(&m); // p1
        if(numfull == max) //p2
            Cond_wait(&cond, &m); //p3
        do_fill(); // p4
        Cond_signal(&cond); //p5
        Mutex_unlock(&m); //p6
    }
}

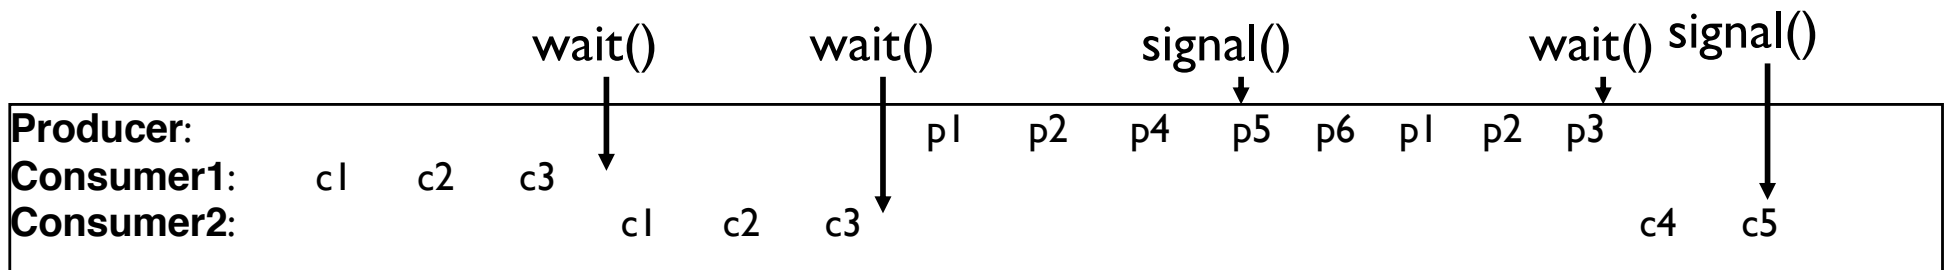
```

```

void *consumer(void *arg) {
    while(1) {
        Mutex_lock(&m); // c1
        if(numfull == 0) // c2
            Cond_wait(&cond, &m); // c3
        int tmp = do_get(); // c4
        Cond_signal(&cond); // c5
        Mutex_unlock(&m); // c6
        do_work(tmp); // c7
    }
}

```

Want consumer2 signal() to wake producer since numbufs = 0, but could wake consumer1  
 Cannot assume which waiting thread will be worken!



# HOW TO WAKE THE RIGHT THREAD?

Wake all the threads!?

# WAKING ALL WAITING THREADS

**wait**(cond\_t \*cv, mutex\_t \*lock)

- assumes the lock is held when wait() is called
- puts caller to sleep + releases the lock (atomically)
- when awoken, reacquires lock before returning

**signal**(cond\_t \*cv)

- wake a single waiting thread (if  $\geq 1$  thread is waiting)
- if there is no waiting thread, just return, doing nothing

**broadcast**(cond\_t \*cv)

- wake **all** waiting threads (if  $\geq 1$  thread is waiting)
- if there are no waiting thread, just return, doing nothing



# HOW TO WAKE THE RIGHT THREAD?

Wake all the threads!?

Better solution (usually): use separate condition variables

# PRODUCER/CONSUMER: TWO CVS

```
void *producer(void *arg) {
    While (1) {
        Mutex_lock(&m); // p1
        if (numfull == max) // p2
            Cond_wait(&empty, &m); // p3
        do_fill(); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}
```

```
void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        if (numfull == 0)
            Cond_wait(&fill, &m);
        int tmp = do_get();
        Cond_signal(&empty);
        Mutex_unlock(&m);
    }
}
```

Is this correct? Can you find a bad schedule?

# PRODUCER/CONSUMER: TWO CVS

```
void *producer(void *arg) {
    while (1) {
        Mutex_lock(&m); // p1
        if (numfull == max) // p2
            Cond_wait(&empty, &m); // p3
        do_fill(); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}
```

```
void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        if (numfull == 0)
            Cond_wait(&fill, &m);
        int tmp = do_get();
        Cond_signal(&empty);
        Mutex_unlock(&m);
    }
}
```

1. consumer1 waits because numfull == 0
2. producer increments numfull, wakes consumer1
3. before consumer1 runs, consumer2 runs, grabs lock, gets data and sets numfull=0.
4. When consumer2 wakes from cond\_wait with lock, reads bad data

# PRODUCER/CONSUMER: TWO CVS

```
void *producer(void *arg) {
    while (1) {
        Mutex_lock(&m); // p1
        if (numfull == max) // p2
            Cond_wait(&empty, &m); // p3
        do_fill(); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}
```

```
void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        if (numfull == 0)
            Cond_wait(&fill, &m);
        int tmp = do_get();
        Cond_signal(&empty);
        Mutex_unlock(&m);
    }
}
```

Cannot assume which threads will acquire lock next

When wake from cond\_wait(), must recheck state to ensure state is indeed true

(i.e., no other thread changed state between cond\_signal() returning from cond\_wait())

# PRODUCER/CONSUMER: TWO CVS AND WHILE

```
void *producer(void *arg) {
    while (1) {
        Mutex_lock(&m); // p1
        while (numfull == max) // p2
            Cond_wait(&empty, &m); // p3
        do_fill(); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}
```

```
void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        while (numfull == 0)
            Cond_wait(&fill, &m);
        int tmp = do_get();
        Cond_signal(&empty);
        Mutex_unlock(&m);
    }
}
```

Lock of mutex: No concurrent access to shared state

Every time lock is re-acquired, assumptions are reevaluated (while loop instead of if)

Progress: A consumer will get to run after every do\_fill()

Progress: A producer will get to run after every do\_get()

# GOOD RULE OF THUMB 3

Whenever a lock is acquired, **recheck assumptions** about state!

Another thread could grab lock in between signal and wakeup from wait

Some implementations have “spurious wakeups”

- May wake multiple waiting threads at signal or at any time
- May treat `signal()` as `broadcast()`
- Good way to stress test your code: change signals to broadcasts

# SUMMARY: RULES OF THUMB FOR CVS

1. Keep state in addition to CV's
2. Always change state and do wait/signal with lock held
3. Whenever thread wakes from waiting, recheck state

# IMPLEMENTING SYNCHRONIZATION

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right

