

# CONCURRENCY: SEMAPHORES

Andrea Arpaci-Dusseau  
CS 537, Fall 019

# ADMINISTRIVIA

- Project 4 due today at 5:00 pm (or 7:00 am...)
- Project 5 available Wed morning (xv6 Memory)
  - Request new project partner if desired
- Midterm 2: Nov 11/6 (Wed) from 7:30-9:30pm
  - Notify by tomorrow if conflict
  - Two "quizzes" on race conditions in Canvas
- Office Hours Changed Today: 4-5 pm

# LEARNING OUTCOMES: SEMAPHORES

**Semaphores** (vs. condition variables?)

How to implement a **lock** with semaphores?

How to implement semaphores with locks and condition variables?

How to implement **join** and **producer/consumer** with semaphores?

How to synchronize **dining philosophers**?

How to implement **reader/writer locks** with semaphores?

**RECAP**

# CONCURRENCY OBJECTIVES

**Mutual exclusion** (e.g., A and B don't run at same time)  
solved with *locks*

**Ordering** (e.g., B runs after A does something)  
solved with *condition variables* and *semaphores*


# CONDITION VARIABLES

**wait**(cond\_t \*cv, mutex\_t \*lock)

- assumes the lock is held when wait() is called
- puts caller to sleep + releases the lock (atomically)
- when awoken, reacquires lock before returning

**signal**(cond\_t \*cv)

- wake a single waiting thread (if  $\geq 1$  thread is waiting)
- if there is no waiting thread, just return, doing nothing



signal(cv) - what happens?

release(lock) - what happens?

# ORDERING EXAMPLE: JOIN

```
pthread_t p1, p2;  
Pthread_create(&p1, NULL, mythread, "A");  
Pthread_create(&p2, NULL, mythread, "B");  
// join waits for the threads to finish (call exit())  
Pthread_join(p1, NULL);  
Pthread_join(p2, NULL);  
printf("main: done\n [balance: %d]\n [should: %d]\n",  
       balance, max*2);  
return 0;
```

how to implement join()?

# JOIN IMPLEMENTATION: CORRECT

Parent:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m);    // y  
    Mutex_unlock(&m);         // z  
}
```

Child:

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1;                // b  
    Cond_signal(&c);          // c  
    Mutex_unlock(&m);         // d  
}
```

Parent: w            x            y                            z

Child:                            a            b            c

Use mutex to ensure no race between interacting with state and wait/signal



# PRODUCER/CONSUMER: TWO CVS AND WHILE

```
void *producer(void *arg) {
    while (1) {
        Mutex_lock(&m); // p1
        while (numfull == max) // p2
            Cond_wait(&empty, &m); // p3
        do_fill(); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}
```

```
void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        while (numfull == 0)
            Cond_wait(&fill, &m);
        int tmp = do_get();
        Cond_signal(&empty);
        Mutex_unlock(&m);
        do_work(tmp);
    }
}
```

# SUMMARY: RULES OF THUMB FOR CVS

1. Keep state in addition to CV's
2. Always do wait/signal with lock held
3. Whenever thread wakes from waiting, recheck state (while loop)

# INTRODUCING SEMAPHORES

Condition variables have no **state** (other than waiting queue)

- Programmer must track additional state

Semaphores have state: **track integer value**

- State cannot be directly accessed by user program,  
but state determines behavior of semaphore operations

# EQUIVALENCE CLAIM

Semaphores are equally powerful to Locks+CVs

- what does this mean?

One might be more convenient, but that's not relevant

Equivalence means each can be built from the other



# SEMAPHORE OPERATIONS

## Allocate and Initialize

```
sem_t sem;  
sem_init(sem_t *s, int initval) {  
    s->value = initval;  
}
```

User cannot read or write value directly after initialization

wait and post are atomic

**Wait or Test (sometime P() for Dutch) **sem\_wait(sem\_t\*)****

Waits until value of sem is  $> 0$ ; Decrements sem value,

**Signal or Post (sometime V() for Dutch) **sem\_post(sem\_t\*)****

Increment sem value, if value  $> 0$ , wake a single waiter

# BUILD LOCK FROM SEMAPHORE

```
typedef struct __lock_t {  
    sem_t sem;  
} lock_t;
```

sem\_init(sem\_t\*, int initial)  
sem\_wait(sem\_t\*): Wait until value > 0; dec  
sem\_post(sem\_t\*): Increment; if > 0, wake single waiter

```
void init(lock_t *lock) {  
  
}
```

```
void acquire(lock_t *lock) {  
  
}
```

```
void release(lock_t *lock) {  
  
}
```

**Locks**

**Semaphores**

# BUILD LOCK FROM SEMAPHORE

```
typedef struct __lock_t {  
    sem_t sem;  
} lock_t;
```

sem\_init(sem\_t\*, int initial)  
sem\_wait(sem\_t\*): Wait until value > 0; dec  
sem\_post(sem\_t\*): Increment; if > 0, wake single waiter

```
void init(lock_t *lock) {  
    sem_init(&lock->sem, ;  
}
```

```
void acquire(lock_t *lock) {  
      
}
```


```
void release(lock_t *lock) {  
      
}
```

**Locks**

**Semaphores**

# BUILDING CV'S OVER SEMAPHORES

Possible, but really hard to do right



Read about Microsoft Research's attempts:

<http://research.microsoft.com/pubs/64242/ImplementingCVs.pdf>



# JOIN WITH CV VS SEMAPHORES

```
void thread_join() {  
    Mutex_lock(&m);      // w  
    if (done == 0)      // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);    // z  
}
```

```
void thread_exit() {  
    Mutex_lock(&m);      // a  
    done = 1;           // b  
    Cond_signal(&c);     // c  
    Mutex_unlock(&m);    // d  
}
```

---

```
sem_t s;  
sem_init(&s, )
```

sem\_wait(sem\_t\*): Wait until value > 0; dec  
sem\_post(sem\_t\*): Increment; if > 0, wake single waiter

```
void thread_join() {  
    sem_wait(&s);  
}
```

```
void thread_exit() {  
    sem_post(&s);  
}
```

# BUILD SEMAPHORE FROM LOCK AND CV

```
Typedef struct {  
    int value;  
    cond_t cond;  
    lock_t lock;  
} sem_t;
```

Sem\_wait(): Waits until value > 0, then decrement  
Sem\_post(): Increment value, then wake a single waiter

```
Void sem_init(sem_t *s, int value) {  
    s->value = value;  
    cond_init(&s->cond);  
    lock_init(&s->lock);  
}
```

Semaphores

Locks

CV's

# BUILD SEMAPHORE FROM LOCK AND CV

```
Sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
    while (s->value <= 0)  
        cond_wait(&s->cond);  
    s->value--;  
    lock_release(&s->lock);  
}
```

```
Sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
    s->value++;  
    cond_signal(&s->cond);  
    lock_release(&s->lock);  
}
```

Sem\_wait(): Waits until value > 0, then decrement  
Sem\_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

# PRODUCER/CONSUMER: SEMAPHORES #1

Single producer thread, single consumer thread

Single shared buffer between producer and consumer

Use 2 semaphores

- emptyBuffer: Initialize to:
- fullBuffer: Initialize to:


Producer

```
while (1) {  
    sem_wait(&emptyBuffer);  
    Fill(&buffer);  
    sem_signal(&fullBuffer);  
}
```

Consumer

```
while (1) {  
    sem_wait(&fullBuffer);  
    Use(&buffer);  
    sem_signal(&emptyBuffer);  
}
```

# PRODUCER/CONSUMER: SEMAPHORES #2

Single producer thread, single consumer thread

Shared buffer with **N elements** between producer and consumer

Use 2 semaphores

- emptyBuffer: Initialize to:
- fullBuffer: Initialize to:


Producer

```
i = 0;
while (1) {
    sem_wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    sem_signal(&fullBuffer);
}
```

Consumer

```
j = 0;
While (1) {
    sem_wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    sem_signal(&emptyBuffer);
}
```

# PRODUCER/CONSUMER: SEMAPHORE #3

Final case:

- Multiple producer threads, multiple consumer threads
- Shared buffer with N elements between producer and consumer

Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element

# PRODUCER/CONSUMER: MULTIPLE THREADS

Does this code work correctly?

Producer

```
while (1) {  
    sem_wait(&emptyBuffer);  
    my_i = findempty(&buffer);  
    Fill(&buffer[my_i]);  
    sem_signal(&fullBuffer);  
}
```

Consumer

```
while (1) {  
    sem_wait(&fullBuffer);  
    my_j = findfull(&buffer);  
    Use(&buffer[my_j]);  
    sem_signal(&emptyBuffer);  
}
```

Are my\_i and my\_j private or shared? Where is mutual exclusion needed???

# PRODUCER/CONSUMER: MULTIPLE THREADS

Consider three possible locations for mutual exclusion  
Which work??? Which is best???

Producer #1

```
sem_wait(&mutex);  
sem_wait(&emptyBuffer);  
my_i = findempty(&buffer);  
Fill(&buffer[my_i]);  
sem_signal(&fullBuffer);  
sem_signal(&mutex);
```

Consumer #1

```
sem_wait(&mutex);  
sem_wait(&fullBuffer);  
my_j = findfull(&buffer);  
Use(&buffer[my_j]);  
sem_signal(&emptyBuffer);  
sem_signal(&mutex);
```



# PRODUCER/CONSUMER: MULTIPLE THREADS

Producer #2

```
sem_wait(&emptyBuffer);  
sem_wait(&mutex);  
myi = findempty(&buffer);  
Fill(&buffer[myi]);  
sem_signal(&mutex);  
sem_signal(&fullBuffer);
```

Consumer #2

```
sem_wait(&fullBuffer);  
sem_wait(&mutex);  
myj = findfull(&buffer);  
Use(&buffer[myj]);  
sem_signal(&mutex);  
sem_signal(&emptyBuffer);
```

Works, but limits concurrency:

Only 1 thread at a time can be using or filling different buffers

# PRODUCER/CONSUMER: MULTIPLE THREADS

Producer #3

```
sem_wait(&emptyBuffer);  
sem_wait(&mutex);  
myi = findempty(&buffer);  
sem_signal(&mutex);  
Fill(&buffer[myi]);  
sem_signal(&fullBuffer);
```

Consumer #3

```
sem_wait(&fullBuffer);  
sem_wait(&mutex);  
myj = findfull(&buffer);  
sem_signal(&mutex);  
Use(&buffer[myj]);  
sem_signal(&emptyBuffer);
```

Works and increases concurrency; only finding a buffer is protected by mutex;  
Filling or Using different buffers can proceed concurrently

# DINING PHILOSOPHERS

## Problem Statement

- **N** Philosophers sitting at a round table
- Each philosopher shares a chopstick (or fork) with neighbor
- Each philosopher must have both chopsticks to eat
- Neighbors can't eat simultaneously
- Philosophers alternate between thinking and eating

Each philosopher/thread **i** runs :

```
while (1) {  
    think();  
    take_chopsticks(i);  
    eat();  
    put_chopsticks(i);  
}
```



# DINING PHILOSOPHERS: ATTEMPT #1

Two neighbors can't use chopstick at same time


Must test if chopstick is there and grab it atomically

Represent each chopstick with a semaphore

Grab right chopstick then left chopstick

Code for 5 philosophers:

```
sem_t chopstick[5]; // Initialize each to 1
take_chopsticks(int i) {
    wait(&chopstick[i]);
    wait(&chopstick[(i+1)%5]);
}
put_chopsticks(int i) {
    signal(&chopstick[i]);
    signal(&chopstick[(i+1)%5]);
}
```



# DINING PHILOSOPHERS: ATTEMPT #1

Two neighbors can't use chopstick at same time

Must test if chopstick is there and grab it atomically


Represent each chopstick with a semaphore

Grab right chopstick then left chopstick

Deadlocked!

Code for 5 philosophers:

```
sem_t chopstick[5]; // Initialize each to 1
take_chopsticks(int i) {
    wait(&chopstick[i]);
    wait(&chopstick[(i+1)%5]);
}
put_chopsticks(int i) {
    signal(&chopstick[i]);
    signal(&chopstick[(i+1)%5]);
}
```



# DINING PHILOSOPHERS: ATTEMPT #2


Grab lower-numbered chopstick first, then higher-numbered

```
sem_t chopstick[5]; // Initialize to 1
take_chopsticks(int i) {
    if (i < 4) {
        wait(&chopstick[i]);
        wait(&chopstick[i+1]);
    } else {
        wait(&chopstick[0]);
        wait(&chopstick[4]);
    }
}
```

Philosopher 3 finishes take\_chopsticks() and eventually calls put\_chopsticks();

Who can run then?

What is wrong with this solution???



# DINING PHILOSOPHERS: HOW TO APPROACH

Guarantee two goals

- **Safety:** Ensure nothing bad happens (don't violate constraints of problem)
- **Liveness:** Ensure something good happens when it can (make as much progress as possible)

Introduce state variable for each philosopher  $i$

`state[i] = THINKING, HUNGRY, or EATING`

## **Safety:**

No two adjacent philosophers eat simultaneously

`for all  $i$ :  $!(state[i]==EATING \ \&\& \ state[i+1\%5]==EATING)$`

## **Liveness:**

Not the case that a philosopher is hungry and his neighbors are not eating

`for all  $i$ :  $!(state[i]==HUNGRY \ \&\& \ (state[i+4\%5]!=EATING \ \&\& \ state[i+1\%5]!=EATING))$`

```

sem_t mayEat[5]; // how to initialize?
sem_t mutex;     // how to init?
int state[5] = {THINKING};
take_chopsticks(int i) {
    wait(&mutex); // enter critical section
    state[i] = HUNGRY;
    testSafetyAndLiveness(i); // check if I can run
    signal(&mutex); // exit critical section
    wait(&mayEat[i]);
}
put_chopsticks(int i) {
    wait(&mutex); // enter critical section
    state[i] = THINKING;
    test(i+1 %5); // check if neighbor can run now
    test(i+4 %5);
    signal(&mutex); // exit critical section
}
testSafetyAndLiveness(int i) {
    if(state[i]==HUNGRY&&state[i+4%5]!=EATING&&state[i+1%5]!=EATING) {
        state[i] = EATING;
        signal(&mayEat[i]);
    }
}

```



# DINING PHILOSOPHERS: EXAMPLE EXECUTION

Take\_chopsticks(0)

Take\_chopsticks(1)


Take\_chopsticks(2)

Take\_chopsticks(3)

Take\_chopsticks(4)

Put\_chopsticks(0)

Put\_chopsticks(2)



```

sem_t mayEat[5]; // how to initialize?
sem_t mutex;     // how to init?
int state[5] = {THINKING};
take_chopsticks(int i) {
    wait(&mutex); // enter critical section
    state[i] = HUNGRY;
    testSafetyAndLiveness(i); // check if I can run
    signal(&mutex); // exit critical section
    wait(&mayEat[i]);
}
put_chopsticks(int i) {
    wait(&mutex); // enter critical section
    state[i] = THINKING;
    test(i+1 %5); // check if neighbor can run now
    test(i+4 %5);
    signal(&mutex); // exit critical section
}
testSafetyAndLiveness(int i) {
    if(state[i]==HUNGRY&&state[i+4%5]!=EATING&&state[i+1%5]!=EATING) {
        state[i] = EATING;
        signal(&mayEat[i]);
    }
}

```

```

Take_chopsticks(0)
Take_chopsticks(1)
Take_chopsticks(2)
Take_chopsticks(3)
Take_chopsticks(4)
Put_chopsticks(0)
Put_chopsticks(2)

```

# READER/WRITER LOCKS

Protect shared data structure; Goal:

Let multiple reader threads grab lock with other readers (shared)

Only one writer thread can grab lock (exclusive)

- No reader threads
- No other writer threads

Two possibilities for priorities – different implementations

1) No reader waits unless writer in critical section

- How can writers starve?

2) No writer waits longer than absolute minimum

- How can readers starve?

Let us see if we can understand code...

# VERSION 1

Readers have priority

# READER/WRITER LOCKS

```
1 typedef struct _rwlock_t {
2     sem_t lock;
3     sem_t writelock;
4     int readers;
5 } rwlock_t;
6
7 void rwlock_init(rwlock_t *rw) {
8     rw->readers = 0;
9     sem_init(&rw->lock, 1);
10    sem_init(&rw->writelock, 1);
11 }
```

# READER/WRITER LOCKS

```
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     rw->readers++;
16     if (rw->readers == 1)
17         sem_wait(&rw->writelock);
18     sem_post(&rw->lock);
19 }
```

```
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0)
25         sem_post(&rw->writelock);
26     sem_post(&rw->lock);
27 }
```

```
29 rwlock_acquire_writelock(rwlock_t *rw) { sem_wait(&rw->writelock); }
31 rwlock_release_writelock(rwlock_t *rw) { sem_post(&rw->writelock); }
```

T1: acquire\_readlock()  
T2: acquire\_readlock()  
T3: acquire\_writelock()  
T2: release\_readlock()  
T1: release\_readlock()  
// who runs?  
T4: acquire\_readlock()  
// what happens?  
T5: acquire\_readlock()  
// where blocked?  
T3: release\_writelock()  
// what happens next?

# READER/WRITER LOCKS

```
13 void rwlock_acquire_readlock(rwlock_t *rw) {
```

```
14     sem_wait(&rw->lock);
```

```
15     rw->readers++;
```

```
16     if (rw->readers == 1)
```

```
17         sem_wait(&rw->writelock);
```

```
18     sem_post(&rw->lock);
```

```
19 }
```

```
21 void rwlock_release_readlock(rwlock_t *rw) {
```

```
22     sem_wait(&rw->lock);
```

```
23     rw->readers--;
```

```
24     if (rw->readers == 0)
```

```
25         sem_post(&rw->writelock);
```

```
26     sem_post(&rw->lock);
```

```
27 }
```

```
29 rwlock_acquire_writelock(rwlock_t *rw) { sem_wait(&rw->writelock); }
```

```
31 rwlock_release_writelock(rwlock_t *rw) { sem_post(&rw->writelock); }
```

T1: acquire\_readlock()

T2: acquire\_readlock()

T3: acquire\_writelock()

T4: acquire\_readlock()

// what happens?

# VERSION 2

Writers have priority

Three semaphores

- Mutex
- OKToRead (siimilar to myEat[] in Dining Philosphers, but 1 for all readers)
- OKToWrite

How to initialize?

Shared variables

Waiting Readers,

ActiveReaders

WaitingWriters

ActiveWriters



```

Acquire_readlock() {
    Sem_wait(&mutex);
    If (ActiveWriters +
        WaitingWriters==0) {
        sem_post(OKToRead);
        ActiveReaders++;
    } else WaitingReaders++;
    Sem_post(&mutex);
    Sem_wait(OKToRead);
}

Release_readlock() {
    Sem_wait(&mutex);
    ActiveReaders--;
    If (ActiveReaders==0 &&
        WaitingWriters > 0) {
        Sem_post(OKToWrite);
        ActiveWriters++;
        WaitingWriters--;
    }
    Sem_post(&mutex);
}

```

```

Acquire_writelock() {
    Sem_wait(&mutex);
    If (ActiveWriters + ActiveReaders + WaitingWriters==0) {
        ActiveWriters++;
        sem_post(OKToWrite);
    } else WaitingWriters++;

    Sem_post(&mutex);
    Sem_wait(OKToWrite);
}

Release_writelock() {
    Sem_wait(&mutex);
    ActiveWriters--;
    If (WaitingWriters > 0) {
        ActiveWriters++;
        WaitingWriters--;
        Sem_post(OKToWrite);
    } else while(WaitingReaders>0) {
        ActiveReaders++;
        WaitingReaders--;
        sem_post(OKToRead);
    }
    Sem_post(&mutex);
}

```

T1: acquire\_readlock()  
 T2: acquire\_readlock()  
 T3: acquire\_writelock()  
 T4: acquire\_readlock()  
 // what happens?

# SEMAPHORES

Semaphores are equivalent to locks + condition variables

- Can be used for both mutual exclusion and ordering

Semaphores contain **state**

- How they are initialized depends on how they will be used
- Init to 0: Join (1 thread must arrive first, then other)
- Init to N: Number of available resources

`Sem_wait()`: Waits until value  $> 0$ , then decrement (atomic)

`Sem_post()`: Increment value, then wake a single waiter (atomic)

Can use semaphores in producer/consumer and for reader/writer locks

# REVIEW: PROCESSES VS THREADS

```
int a = 0;
int main() {
    fork();
    a++;
    fork();
    a++;
    if (fork() == 0) {
        printf("Hello!\n");
    } else {
        printf("Goodbye!\n");
    }
    a++;
    printf("a is %d\n", a);
}
```

How many times will “Hello!\n” be displayed?

4

What will be the **final** value of “a” as displayed by the final line of the program?

3